Distributed consensus algorithms ensure node clusters reach agreement despite network partitions.
Raft simplifies consensus by decomposing it into leader election, log replication, and safety guarantees.
A leader is elected via randomized heartbeat timers to prevent split-vote scenarios.
Once established, all client state transitions flow strictly through the leader before being committed to follo
Append-only logs guarantee consistent state machine replication across all participating replicas.